



Advanced Python for Neuroscientists

Lecture 8: Reinforcement learning

Summer 2022

Princeton Neuroscience Institute
Instructors: Yisi Zhang & Bichan Wu

July 21, 2022



Recap

Lecture 5: Neural network I

- Feedforward neural network
- Gradient descent

Lecture 6: Neural network II

- Backpropagation
- Stochastic Gradient Descent
- Application

Lecture 7: Convolutional neural network

- Motivation & concept
- Overview
- Architectures





Outline

- Motivation & concept
- Formalization
- Deep Q-learning



8.1 Motivation & concept

How do you train a machine to play games?



How to train this model using supervised learning?
In this case, what is the input and what is the output?



8.1 Motivation & concept

Issues with SL:


- Too many possible solutions (ground truth not well defined)
- Cost function not well defined
- Cannot generalize

Why RL?

- Learn from interaction
- Does not require a knowledgeable external supervisor
- Make sequence of decisions

8.1 Motivation & concept

Reinforcement learning




An **agent** interacting with an **environment**, which provides numeric **reward**.

Goal: Learn how to take actions to maximize reward



8.1 Motivation & concept

Cart-pole problem



Objective: Balance a pole on top of a moving cart

State: Angle, angular speed, position, horizontal velocity


Action: Horizontal force applied to the cart

Reward: 1 at each time step if the pole is upright



8.1 Motivation & concept

Go



Objective: Win the game

State: Position of all pieces

Action: Where to put the next piece

Reward: 1 if win

8.2 Formalization

Markov decision process

- Markov: the environment's response at $t + 1$ depends only on the state and action representations at t

Transition probability

$$\mathcal{P}_{ss'}^a = Pr\{s_{t+1} = s' | s_t = s, a_t = a\}$$

Expected reward

$$\mathcal{R}_{ss'}^a = E\{r_{t+1} | s_t = s, a_t = a, s_{t+1} = s'\}$$

8.2 Formalization

Markov decision process

- At time step $t = 0$, environment samples initial state $s_0 \sim p(s_0)$
- Then, for $t = 0$ until done:
 - Agent selects action a_t
 - Environment samples reward $r_t \sim R(.|s_t, a_t)$
 - Environment samples next state $s_{t+1} \sim P(.|s_t, a_t)$
 - Agent receives reward r_t and next states s_{t+1}
- A policy π is a function from S to A that specifies what action to take in each state
- Objective: find policy π^* that maximizes the **expected return**



8.2 Formalization

Expected return

$$R_t = r_{t+1} + \gamma r_{t+2} + \gamma^2 r_{t+3} + \dots = \sum_k \gamma^k r_{t+k+1}$$

where $0 \leq \gamma \leq 1$, called discount rate.

8.2 Formalization

Example: Grid world

states

	1	2	3
4	5	6	7
8	9	10	11
12	13	14	


actions = {
1. right →
2. left ←
3. up ↑
4. down ↓
}

Objective: reach the terminal grids in least number of actions


8.2 Formalization

Example: Grid world

Random policy



Optimal policy





8.2 Formalization

The optimal policy π^*

We want to find the **optimal policy** to maximize the expected return (sum of rewards).

$$\pi^* = \operatorname{argmax}_{\pi} \mathbb{E} \left[\sum_{t \geq 0} \gamma^t r_t \mid \pi \right]$$

with $s_0 \sim p(s_0)$, $a_t \sim \pi(\cdot | s_t)$, $s_{t+1} \sim p(\cdot | s_t, a_t)$



8.2 Formalization

Value functions

To evaluate how good a state is:

State-value function for policy π

$$V^\pi(s) = E_\pi\{R_t | s_t = s\} = E_\pi\left\{\sum_{k=0}^{\infty} \gamma^k r_{t+k+1} | s_t = s\right\}$$



8.2 Formalization

Value functions

To evaluate how good an action is in state s :

Action-value function for policy π

$$Q^{\pi}(s, a) = E_{\pi}\{R_t | s_t = s, a_t = a\} = E_{\pi}\left\{\sum_{k=0}^{\infty} \gamma^k r_{t+k+1} | s_t = s, a_t = a\right\}$$



8.2 Formalization

Bellman equation

The optimal Q-value function Q^* is one which yields maximum expected return achievable from (s, a) :

$$Q^*(s, a) = E\{r_{t+1} + \gamma \max_{a'} Q^*(s', a') | s_t = s, a_t = a\}$$

8.2 Formalization

Q-learning

Temporal difference (TD)

$$Q_t(s, a) = Q_{t-1}(s, a) + \alpha \underbrace{(R(s, a) + \gamma \max_{a'} Q(s', a') - Q_{t-1}(s, a))}_{\text{TD}}$$



8.2 Formalization

Grid world: find the shortest path to terminals (coding time)

Q-table

	up	down	left	right
0				
1				
2				
3				
4				
5				
6				
7				
8				
9				
10				
11				
12				
13				
14				
15				



8.3 Deep Q-learning

Use a function approximator to estimate the action-value function.

$$Q(s, a; \theta) \approx Q^*(s, a)$$

If the approximator is a deep neural network, it is called **Deep Q-learning**

8.3 Deep Q-learning

We want to find

$$Q^*(s, a) = E\{r + \gamma \max_{a'} Q^*(s', a') | s, a\}$$

Forward pass

$$\text{Cost function } J_i(\theta_i) = [y_i - Q(s, a; \theta_i)]^2$$

$$\text{where } y_i = E\{r + \gamma \max_{a'} Q^*(s', a'; \theta_{i-1}) | s, a\}$$

Backward pass


Gradient descent with respect to Q-function parameters θ

$$\nabla_{\theta_i} J_i(\theta_i) =$$

$$(E\{r + \gamma \max_{a'} Q^*(s', a'; \theta_{i-1})\} - Q(s, a; \theta_i)) \nabla_{\theta_i} Q(s, a; \theta_i)$$

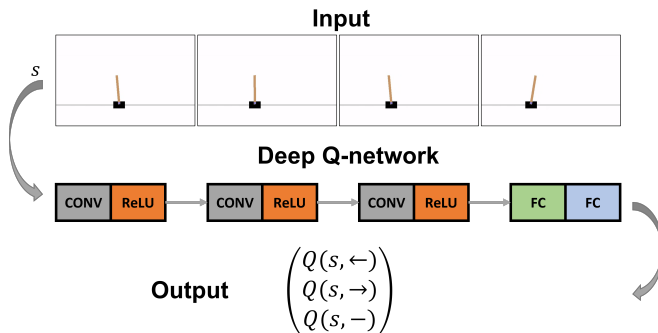
8.3 Deep Q-learning

Implement neural network



8.3 Deep Q-learning

Example: cartpole (coding time)





Homework

- Make sure you understand all the exercises above
- Run through the codes here that should replicate all the figures
<https://github.com/yisiszhang/AdvancedPython/blob/main/colab/Lecture8.ipynb>
- Try to improve the Q-learning code